

Market Deformation Geometry

A Teichmüller Framework for Structural Change Detection
in Multi-Asset, Multi-Venue Markets

Marc Brendecke

Independent Research

September 2026

Abstract

This paper proposes *Market Deformation Geometry*: a structure-first Teichmüller framework that measures structural change in multi-asset, multi-venue financial markets as quasiconformal deformation of synchronized price observations. The observable market is represented not as a collection of isolated price charts, but as a coupled multidimensional geometric object. Foreign exchange serves as the running example throughout, because it supplies many algebraically related instruments and multiple observable price feeds; the construction itself applies to any universe of instruments observed through several venues.

The fundamental data structure is the synchronized tensor $P(t, a, b)$, where t denotes time, a an asset — a currency pair, a metal, an index, or any other instrument — and b a broker or price-feed observation. The temporal coordinate describes the evolution of the market, the asset coordinate describes cross-asset relationships, and the broker coordinate describes different realizations of nominally identical market instruments.

The central observation is that these dimensions must not be treated as interchangeable. Connecting observations along the time axis produces a conventional price path. Connecting observations along the broker axis instead produces a cross-sectional observation geometry. Connecting assets produces a third structural direction. The paper develops finite-difference analogues of $\partial_t P$, $\partial_a P$, $\partial_b P$ together with the mixed derivatives $\partial_t \partial_a P$, $\partial_t \partial_b P$, $\partial_a \partial_b P$. Particular attention is given to $\partial_t \partial_b P$, which measures how the temporal evolution of an instrument differs across broker observations.

The second part of the framework introduces a quasiconformal geometric layer. Local market patches are normalized and compared by deformation maps; the Beltrami coefficient $\mu = f_{\bar{z}}/f_z$ and the quasiconformal dilatation $K = (1 + |\mu|)/(1 - |\mu|)$ provide candidate measures of structural deformation. The proposed financial interpretation is deliberately conservative: the market is not assumed to literally be a Riemann surface. Instead, market data are embedded into finite-dimensional geometric patches, and quasiconformal mappings are used as a mathematical comparison device.

The literature already contains geometric approaches to finance — manifold learning, Riemannian and Minkowski formulations, topological and cohomological constructions, and a direct use of Teichmüller and moduli-space ideas in financial time-series classification. The proposed contribution is therefore not “geometry in finance” in general; the specific gap addressed is the *combination* of synchronized multi-asset observations, multiple observation venues, explicit quasiconformal deformation, and a Teichmüller-type market distance.

The resulting framework defines a market state in terms of geometric deformation, cross-asset coherence, cross-broker coherence, mixed-dimensional coupling, and triangular consistency, and formulates empirical hypotheses, null models, estimator validation, falsification tests, and potential applications to regime detection. A statistical-power pre-calculation determines which hypotheses the planned dataset can decide: window-level coherence hypotheses within the first cycle, event-level transition hypotheses only for large effects, and economic significance not at all on realistic horizons. A first implementation cycle sharpens both the mathematics and the empirics: on simply connected patches the proposed distance is identically zero by the Riemann mapping theorem, and the tractable carrier of Teichmüller

structure is the torus, where the distance acquires a closed form on conformal classes of 2×2 covariances and coordinate robustness (H7) becomes a theorem for that estimator; empirically, the venue axis fails as a geometric direction but serves as a replication axis, giving the conformal shape a measured error floor roughly a factor of ten below its movement through time — though against the sampling noise of the window covariance, rather than against venue noise, the same movement is only a factor of two above the floor, and a preregistered volatility-clustering null narrows but does not close the gap. The primary research question is not whether a new indicator can be invented, but whether the multidimensional geometry of synchronized market observations contains information that cannot be reduced to volatility, correlation, PCA, cointegration, or standard technical indicators.

Contents

1	Introduction	3
2	Related Work and the Research Gap	3
3	The Market Tensor $P(t, a, b)$	4
4	Temporal Synchronization of Multi-Venue Data	5
5	Logarithmic Price Coordinates	5
6	Currency-Network Representation	5
7	The Currency Graph and Its Incidence Matrix	6
8	Triangular No-Arbitrage Consistency	6
9	The Three Geometric Directions: Time, Asset, Broker	6
9.1	Temporal direction	6
9.2	Asset direction	6
9.3	Broker direction	7
10	The Broker Fiber: Cross-Venue Price Dispersion	7
11	Mixed Temporal–Broker Geometry	7
12	Cross-Asset Temporal Geometry	8
13	First- and Second-Order Differential Structure	8
14	Local Market Surfaces and the Induced Metric	8
15	The Tensor as a Family of Coupled Surfaces	8
16	Quasiconformal Maps and the Beltrami Coefficient	9
17	Financial Interpretation of the Beltrami Field	9
18	A Teichmüller-Type Deformation Distance	9
18.1	Degeneracy on simply connected patches	10
18.2	The torus: an exact, closed-form carrier	10
19	Market Teichmüller Space	10

20 Price Movement versus Structural Movement	11
21 Cross-Broker Coherence	11
22 Cross-Asset Coherence	11
23 Joint Market Coherence	11
24 The Geometric Market State Vector	11
25 The Coherence–Deformation Phase Diagram	12
26 A Geometry-Driven Dynamic Price Channel	12
27 Beltrami Phase and Directional Asymmetry	12
28 The Three Deformation Families	12
28.1 Temporal deformation	13
28.2 Asset deformation	13
28.3 Broker deformation	13
29 Mixed Deformation Statistics	13
30 A Discrete Quasiconformal Approximation via the Jacobian	13
31 Estimating a Financial Beltrami Field	13
32 Surface Geometry and the Metric Tensor	14
33 Market Geometry as a Fibered Structure	14
34 Recovering a Latent Common Market Geometry	14
35 Broker Observation Operators	15
36 Cross-Broker Consensus Geometry	15
37 Cross-Asset Consensus Geometry	15
38 Coordinated and Fragmented Regimes	16
39 Research Hypotheses	16
40 Comparison with Correlation, PCA, Cointegration, and Optimal Transport	16
40.1 Correlation	16
40.2 Principal component analysis	17
40.3 Cointegration	17
40.4 Optimal transport	17
40.5 Manifold learning	17
41 Null Models and Estimator Validation	17
41.1 Broker permutation	17
41.2 Time shuffle	17
41.3 Asset shuffle	17
41.4 Synthetic models	17

41.5 Synthetic deformation recovery	18
41.6 Coordinate-invariance test	18
42 Empirical Dataset: FX Core and Cross-Asset Extension	18
43 Statistical Power: What the Dataset Can Decide	18
43.1 Three classes of hypotheses, three limiting sample units	18
43.2 Direction of the approximations	19
44 Rolling Window Estimation	20
45 Data Leakage and Look-Ahead Bias	20
46 Trading Interpretation: From Geometry to Regime	20
47 Transaction Costs and Broker Reality	20
48 Feed Anomaly Detection	21
49 Market Regime Classification	21
50 Multi-Timeframe Extension	21
51 Order-Book and Liquidity-Depth Extension	21
52 Generalization Beyond FX	21
53 The Geometric Hierarchy: From Price Chart to Market Tensor	22
54 Research Programme	22
55 Falsification Criteria	23
56 Implementation Architecture	23
57 Minimal First Experiment	24
58 First Empirical Cycle: Implementation Results	24
58.1 Stage I: the tensor is constructible, and the venue axis has measurable admission costs	25
58.2 Stage II: the venue axis fails as a geometric direction — and succeeds as a replication axis	25
58.3 The measurement floor of the conformal shape	25
58.4 The estimation floor: what the quotient of 9.8 is measured against	26
59 Expected Contributions	28
59.1 A multidimensional market representation	28
59.2 Mixed-dimensional market statistics	28
59.3 Geometric regime measurement	28
59.4 Observation-consensus geometry	28
60 Limitations	28
61 Conclusion	29

A	Notation	29
B	Core Equations	30
C	Recommended Empirical Benchmark Matrix	31
D	Final Research Criterion	31
E	Research-Status Matrix	31
F	Open Mathematical Questions	32

1 Introduction

Conventional financial chart analysis begins with a one-dimensional time series $P(t)$. The independent variable is time and the dependent variable is price. From this representation one constructs moving averages, oscillators, volatility estimators, channels, trend lines, and other derived statistics.

This representation is useful, but it removes information that is intrinsic to real markets. Foreign exchange makes the loss especially visible and serves as the running example; the construction is not restricted to it (Section 52).

A currency pair is not an isolated object. It is related to other currency pairs through currency identities, triangular relationships, liquidity conditions, funding relationships, order flow, and market microstructure. A metal, an index, or a crypto-asset is likewise coupled to the rest of the universe — through common risk factors, funding currencies, and shared liquidity — even where no exact algebraic identity exists.

Furthermore, the same nominal instrument can be observed through multiple brokers or liquidity sources. These observations are similar but generally not identical.

Consequently, instead of considering only $P(t)$, we propose the multidimensional object

$$\boxed{P(t, a, b)}$$

with

$$t : \text{time}, \quad a : \text{asset / instrument}, \quad b : \text{broker / observation source}.$$

This creates three distinct directions. For fixed (a, b) , the map $t \mapsto P(t, a, b)$ is the conventional price chart. For fixed (t, a) , the map $b \mapsto P(t, a, b)$ describes the cross-broker realization of the same instrument at the same moment. For fixed (t, b) , the map $a \mapsto P(t, a, b)$ describes the cross-asset market state.

The complete object therefore contains substantially more structure than an ordinary chart. The main proposition of this paper is:

The geometry of the multidimensional observation structure may contain information about market regimes that is not visible in an individual price path.

The proposed framework is deliberately structure-first. No trading rule is assumed in advance. The research sequence is

market data $\rightarrow$ geometric representation $\rightarrow$ deformation $\rightarrow$ market state $\rightarrow$ empirical validation.

2 Related Work and the Research Gap

Geometric approaches to finance are established, and the present framework must be positioned against them rather than beside them. Differential-geometric methods have long been used in asset pricing and incomplete-market models, while more recent work embeds financial observations into low-dimensional manifolds. A directly relevant recent example is the use of constant-curvature two-manifolds as latent geometries for financial data, with spherical, Euclidean, hyperbolic, and toroidal candidates compared empirically [20]. Pinčák and Kanjamapornkul explicitly couple price and time through a Minkowski metric in a GARCH framework [21], and a cohomological and topological treatment of financial time series — including link and knot structures and a topological interpretation of bid-ask dynamics — has also been proposed [15].

The closest direct prior art identified is the *Support Spinor Machine* [14], which uses moduli-space and Teichmüller-space ideas as part of a nonlinear time-series classification architecture.

It would therefore be incorrect to claim that Teichmüller concepts have never been applied to financial time series. The claim of this paper is narrower. The intended research gap is

synchronized multi-asset observations $\times$ multiple venues $\times$ explicit quasiconformal deformation $\times$ Teichmüller-type market distance
--

— that is, the specific combination, not any individual ingredient.

Two further strands border the empirical part of this paper, and the boundary should be drawn precisely. First, multi-venue econometrics treats the dispersion between venues as a nuisance: fragmentation noise to be modeled and removed so that integrated variance can be estimated consistently [6]. The present framework inverts that stance — venue dispersion is retained as a replicate-based error bar for geometric statistics (Section 58); the analytic null distribution for geodesic distances between correlation matrices [16], built on the quotient-affine geometry of [23], is the model-based complement to such a measured floor, not a substitute for it. Second, the slow evolution of market correlation structure is an established econophysics finding: correlation-matrix similarity defines persistent market states [19], and the correlation structure of U.S. industry indices evolves on well-separated time scales — the average correlation on a scale beyond five years, the leading spectral quantities in event-driven bursts [5]. That phenomenon is not claimed as new here; what this paper adds on that front is the instrument — a scale- and $GL(2)$ -invariant distance with a closed form and a proof of extremality, reported together with its measured error floor.

Classical quasiconformal and Teichmüller theory provides the mathematical foundation; the standard modern framework is Gardiner and Lakic [13], with [1] and [2] as classical references. In local complex coordinates, a quasiconformal map satisfies

$$f_{\bar{z}} = \mu(z) f_z, \quad \|\mu\|_\infty < 1, \quad (2.1)$$

with the local distortion $K(z)$ developed in Section 16.

To keep the boundary between foundation and proposal visible, every construction in this paper is assigned one of three statuses — *established mathematics*, *proposed construction*, or *research hypothesis* — and Appendix E records this assignment per component.

3 The Market Tensor $P(t, a, b)$

Let $\mathcal{T} = \{t_1, \dots, t_T\}$ be a common time grid, let $\mathcal{A} = \{a_1, \dots, a_N\}$ be the set of instruments — in the FX example currency pairs, in general any mix of asset classes — and let $\mathcal{B} = \{b_1, \dots, b_M\}$ be the set of brokers or independent price-feed sources.

The synchronized price tensor is defined as

$$P_{ijk} = P(t_i, a_j, b_k). \quad (3.1)$$

Thus the complete data structure has $T \times N \times M$ price observations before considering bid, ask, OHLC, volume, spread, or order-book information.

Definition 3.1 (Market Tensor). *The synchronized market tensor is*

$$\mathcal{X} = \{P(t, a, b) \mid t \in \mathcal{T}, a \in \mathcal{A}, b \in \mathcal{B}\}.$$

The tensor should be interpreted as an observation structure rather than as a conventional three-dimensional physical space. The coordinates have different meanings; (t, a, b) is therefore *not* symmetric under arbitrary coordinate permutations.

4 Temporal Synchronization of Multi-Venue Data

Cross-broker comparison requires a common temporal coordinate. Suppose broker b_1 constructs a five-minute bar over $[12:00, 12:05)$ while broker b_2 constructs a bar over $[12:01, 12:06)$. The resulting bars cannot be compared directly without introducing an artificial phase displacement.

We therefore define a canonical partition $\mathcal{I} = \{I_1, \dots, I_T\}$ and a synchronization operator

$$\mathcal{S} : P^{(b)} \mapsto P_{\mathcal{I}}^{(b)}.$$

All brokers must be mapped onto the same canonical intervals.

For tick data, the synchronization problem is more subtle because quotes arrive asynchronously. Possible approaches include:

1. previous-tick interpolation;
2. linear interpolation;
3. event-time sampling;
4. nearest-neighbor synchronization;
5. volume-clock or transaction-clock synchronization.

The chosen procedure must be fixed before statistical testing.

5 Logarithmic Price Coordinates

Prices are multiplicative quantities. Define

$$x(t, a, b) = \log P(t, a, b). \quad (5.1)$$

The temporal return becomes

$$r_t(t, a, b) = x(t + \Delta t, a, b) - x(t, a, b). \quad (5.2)$$

Logarithmic coordinates are particularly useful because cross-rate relationships become additive. For example, $\text{EUR/JPY} = (\text{EUR/USD})(\text{USD/JPY})$ implies

$$\log(\text{EUR/JPY}) = \log(\text{EUR/USD}) + \log(\text{USD/JPY}).$$

This provides a natural connection between the asset dimension and the underlying currency dimension.

6 Currency-Network Representation

This section and the two that follow exploit relational structure that is specific to currencies. In a broader universe they apply to the FX subgraph; instruments without exact algebraic cross-identities simply contribute no closure constraints, while the temporal and broker dimensions apply to them unchanged.

Let $\mathcal{C} = \{c_1, \dots, c_K\}$ denote the currency universe. An FX pair i/j can be represented through latent currency coordinates $u_i(t, b)$ and $u_j(t, b)$. Define

$$x_{ij}(t, b) = u_i(t, b) - u_j(t, b) + \varepsilon_{ij}(t, b). \quad (6.1)$$

The residual ε_{ij} contains quotation noise, microstructure effects, timing effects, and deviations from the latent currency representation.

The representation has a gauge freedom because $(u_i + \lambda) - (u_j + \lambda) = u_i - u_j$. We therefore impose

$$\sum_{i=1}^K u_i(t, b) = 0. \quad (6.2)$$

This converts the currency system into a normalized coordinate system.

7 The Currency Graph and Its Incidence Matrix

Let $G_A = (V_A, E_A)$ be the directed currency graph. Currencies form vertices, $V_A = \mathcal{C}$; FX pairs form directed edges, $E_A \subseteq \mathcal{C} \times \mathcal{C}$.

Let B denote the incidence matrix. Then the vector of pairwise log prices may be approximated by

$$\mathbf{x}(t, b) = B\mathbf{u}(t, b) + \boldsymbol{\varepsilon}(t, b). \quad (7.1)$$

This representation is useful because it allows the cross-asset layer to be interpreted as a projection of a smaller currency state.

8 Triangular No-Arbitrage Consistency

For three currencies i, j, k , the identity $x_{ij} + x_{jk} + x_{ki} = 0$ holds in the idealized frictionless case. Define the triangular residual

$$\tau_{ijk}(t, b) = x_{ij}(t, b) + x_{jk}(t, b) + x_{ki}(t, b). \quad (8.1)$$

In real markets $\tau_{ijk} \neq 0$ because of spreads, transaction costs, latency, quotation conventions, and liquidity constraints. The residual should therefore be treated as a structural observable, not as an automatic arbitrage signal.

Exact closure identities of this kind are a privilege of the currency subgraph. Other asset classes contribute analogous but weaker consistency relations — the spot–futures basis of a metal or index, or the cross-venue identity of the same crypto-asset — and where no such relation exists, the venue dimension alone supplies the consistency structure: the same instrument observed twice must agree up to the broker operators of Section 35.

9 The Three Geometric Directions: Time, Asset, Broker

The market tensor possesses three fundamental directional variations.

9.1 Temporal direction

The discrete temporal derivative is

$$\Delta_t P_{ijk} = P_{i+1,j,k} - P_{i,j,k}. \quad (9.1)$$

This is the ordinary direction of chart movement.

9.2 Asset direction

For appropriately normalized asset coordinates,

$$\Delta_a P_{ijk} = P_{i,j+1,k} - P_{i,j,k}. \quad (9.2)$$

This represents cross-asset structural variation. Because arbitrary numerical ordering of assets is generally meaningless, a practical implementation should preferably replace the raw index j with a graph, factor, or manifold coordinate.

9.3 Broker direction

For fixed time and asset,

$$\Delta_b P_{ijk} = P_{i,j,k+1} - P_{i,j,k}. \quad (9.3)$$

This is the crucial observation dimension. It is not another time series. It represents the dispersion of nominally identical market observations. The ordering caveat of the asset direction applies here with full force: venue indices carry no canonical order, so any statistic built on $\Delta_b P$ must be symmetrized — in practice through pairwise differences (Section 58).

10 The Broker Fiber: Cross-Venue Price Dispersion

For fixed (t, a) define the broker fiber

$$\mathcal{F}_{t,a} = \{P(t, a, b) : b \in \mathcal{B}\}. \quad (10.1)$$

A robust central observation is

$$\bar{x}(t, a) = \text{median}_b x(t, a, b). \quad (10.2)$$

The broker deviation is

$$d_b(t, a, b) = x(t, a, b) - \bar{x}(t, a). \quad (10.3)$$

A robust dispersion measure is

$$\sigma_b(t, a) = 1.4826 \text{ median}_b |d_b(t, a, b)|. \quad (10.4)$$

The coefficient 1.4826 converts the MAD to the standard-deviation scale under Gaussian assumptions.

11 Mixed Temporal–Broker Geometry

The most important broker statistic is not simply $\Delta_b P$. Instead, consider its temporal evolution. Define

$$\Delta_t \Delta_b P_{ijk} = (P_{i+1,j,k+1} - P_{i+1,j,k}) - (P_{i,j,k+1} - P_{i,j,k}). \quad (11.1)$$

In continuous notation, $\partial_t \partial_b P$. This quantity answers:

Do two broker observations evolve differently during the same market movement?

A small broker spread that remains structurally stable may be uninteresting. A rapidly changing broker differential may indicate:

- liquidity fragmentation;
- feed latency;
- different liquidity providers;
- spread expansion;
- asynchronous quote updates;
- market microstructure stress;
- timestamp problems.

12 Cross-Asset Temporal Geometry

Analogously, $\Delta_t \Delta_a P$ measures how temporal movement differs across assets. This is fundamentally different from simple correlation.

Correlation asks: *do two return series move together?* The mixed derivative asks: *how does the local temporal geometry vary across the asset dimension?*

The distinction becomes important during transitions. Two assets may retain high correlation while their local geometric structure changes dramatically.

13 First- and Second-Order Differential Structure

The differential of the market field is formally

$$dP = P_t dt + P_a da + P_b db. \quad (13.1)$$

The second-order differential contains

$$\begin{aligned} d^2 P &= P_{tt} dt^2 + P_{aa} da^2 + P_{bb} db^2 \\ &+ 2P_{ta} dt da + 2P_{tb} dt db + 2P_{ab} da db. \end{aligned} \quad (13.2)$$

The six second-order components have different interpretations. P_{tt} describes temporal curvature; P_{aa} describes curvature across assets; P_{bb} describes curvature across broker observations. The mixed components P_{ta} , P_{tb} , P_{ab} describe interactions between the three market dimensions.

14 Local Market Surfaces and the Induced Metric

Fix an asset a and define a broker-time surface

$$S_a(t, b) = x(t, a, b). \quad (14.1)$$

The local tangent vectors are $S_t = \partial_t S$ and $S_b = \partial_b S$. The induced metric is

$$g = \begin{pmatrix} \langle S_t, S_t \rangle & \langle S_t, S_b \rangle \\ \langle S_b, S_t \rangle & \langle S_b, S_b \rangle \end{pmatrix}. \quad (14.2)$$

Define

$$\rho_{tb} = \frac{g_{tb}}{\sqrt{g_{tt}g_{bb}}}. \quad (14.3)$$


This is a normalized local coupling between temporal and broker variation. Similarly, for fixed broker b , the surface $S_b(t, a) = x(t, a, b)$ produces an asset-time surface, and

$$\rho_{ta} = \frac{g_{ta}}{\sqrt{g_{tt}g_{aa}}} \quad (14.4)$$

measures temporal-asset coupling.

15 The Tensor as a Family of Coupled Surfaces

The three surfaces (t, a) , (t, b) , and (a, b) are projections of the same underlying tensor.



The complete market geometry should therefore be regarded as a family of coupled surfaces rather than as a single chart.

16 Quasiconformal Maps and the Beltrami Coefficient

A central question is how one local market geometry changes into another. Let Φ_0 be a reference market patch and Φ_1 a later market patch. We seek a map

$$f : \Phi_0 \rightarrow \Phi_1. \quad (16.1)$$

For a complex coordinate $z = x + iy$, a sufficiently regular map can be written locally as $df = f_z dz + f_{\bar{z}} d\bar{z}$. The Beltrami coefficient is

$$\mu(z) = \frac{f_{\bar{z}}(z)}{f_z(z)}. \quad (16.2)$$

For an orientation-preserving quasiconformal map, $|\mu(z)| < 1$. The local dilatation is

$$K(z) = \frac{1 + |\mu(z)|}{1 - |\mu(z)|}. \quad (16.3)$$

When $\mu = 0$ the transformation is locally conformal and $K = 1$. As $|\mu| \rightarrow 1$, the distortion increases.

17 Financial Interpretation of the Beltrami Field

The proposed interpretation is

$$\boxed{\mu \quad \leftrightarrow \quad \text{local market deformation}}$$

and

$$\boxed{K \quad \leftrightarrow \quad \text{magnitude of structural distortion.}}$$

This interpretation must be treated as a modeling construction. It is *not* claimed that the observed market is intrinsically a complex manifold. Instead, the financial data are embedded into a local two-dimensional coordinate patch.

One possible coordinate system is $z = u + iv$ with u a normalized time and v a normalized cross-sectional coordinate. The cross-sectional coordinate may represent:

1. broker rank;
2. a latent currency factor;
3. a PCA coordinate;
4. a diffusion-map coordinate;
5. a graph coordinate.

18 A Teichmüller-Type Deformation Distance

Let Φ_0, Φ_1 be normalized market patches. Define the minimal quasiconformal distortion

$$K_{\min} = \inf_f \sup_z K_f(z), \quad (18.1)$$

where f ranges over the admissible deformation maps. We then define the proposed geometric distance

$$D_T(\Phi_0, \Phi_1) = \frac{1}{2} \log K_{\min}. \quad (18.2)$$

This is inspired by the standard logarithmic quasiconformal metric structure of Teichmüller theory. The financial quantity should therefore be interpreted as a *Teichmüller-inspired deformation statistic*.

18.1 Degeneracy on simply connected patches

The definition above has a failure mode that is easy to miss and fatal if missed. If two market patches are simply connected plane domains without marked points, the Riemann mapping theorem supplies a *conformal* map between them, hence $K_{\min} = 1$ and $D_T \equiv 0$ — for every pair of market states. Whatever a general-purpose implementation reports as nonzero on such patches is a property of its own normalization, not of the market. A nontrivial Teichmüller distance requires actual moduli: marked points, or nontrivial topology.

18.2 The torus: an exact, closed-form carrier

Exactly one tractable carrier is available in this setting. A flat metric on the torus is a positive-definite 2×2 matrix; its conformal class is that matrix modulo a positive scale factor; and $\text{SPD}(2)/\mathbb{R}_+$ is the Teichmüller space of the torus. The extremal map between flat tori is affine (Teichmüller's theorem; the quadratic differential is dz^2), so the inf-sup in (18.1) has a closed solution. With Σ_1, Σ_2 the shape matrices normalized to determinant one, $\Lambda = \Sigma_1^{-1}\Sigma_2$, $T = \text{tr } \Lambda$, and $\lambda_+ = (T + \sqrt{T^2 - 4})/2$,

$$d_T = \frac{1}{2} \log \lambda_+. \quad (18.3)$$

Up to a constant factor this is the affine-invariant metric on covariance matrices. The Teichmüller provenance does not supply a new number — it supplies the *proof that this number is the extremal one*, and it removes the optimization that Stage IV of the research programme otherwise requires on two-dimensional patches.

What (18.3) buys beyond a correlation is precisely one property: d_T is invariant under $\text{GL}(2)$ changes of basis of the plane (modulo scale), whereas the PCA concentration $\lambda_1/\text{tr } \Sigma$ is not. For this estimator, coordinate robustness (H7) is therefore a theorem rather than a hypothesis; the hypothesis remains open for every other estimator in this paper. Two boundaries stay in force: $\text{GL}(2)$ -invariance protects against coordinates *within* the plane, not against the choice of the plane itself, which must be preregistered; and the identification is exact only in dimension two — for larger universes the affine-invariant geometry of $\text{SPD}(N)$ remains available, but should then be called by that name.

19 Market Teichmüller Space

Let $\mathfrak{M}$ denote the set of normalized market patches. Introduce an equivalence relation $\Phi_1 \sim \Phi_2$ when the two patches differ only through the chosen admissible normalizations and conformal transformations. The resulting quotient is denoted

$$\mathcal{T}_{\text{mkt}} = \mathfrak{M}/\sim. \quad (19.1)$$

A market state becomes $[\Phi_t] \in \mathcal{T}_{\text{mkt}}$, and a market trajectory is $t \mapsto [\Phi_t]$. The corresponding geometric velocity is

$$V_G(t) = \frac{D_T(\Phi_t, \Phi_{t+\Delta t})}{\Delta t}. \quad (19.2)$$

The geometric acceleration is

$$A_G(t) = V_G(t) - V_G(t - \Delta t). \quad (19.3)$$

20 Price Movement versus Structural Movement

A geometrically stable market regime satisfies approximately $D_T(\Phi_t, \Phi_{t+\Delta t}) \approx 0$. A transition satisfies $D_T(\Phi_t, \Phi_{t+\Delta t}) \uparrow$. A rapidly developing structural transition satisfies $A_G(t) \gg 0$.

This gives a conceptual distinction between price movement and structural movement. A market may exhibit $|\Delta P| \gg 0$ while $D_T \approx 0$. This would correspond to a strong movement that preserves the local market geometry.

Conversely, $|\Delta P| \approx 0$ may coexist with $D_T \gg 0$. This would indicate a structural transition occurring before a large-amplitude price movement.

21 Cross-Broker Coherence

Define broker return dispersion

$$\sigma_B^2(t, a) = \text{Var}_b [r_t(t, a, b)]. \quad (21.1)$$

A normalized coherence statistic may be defined as

$$C_B(t, a) = 1 - \frac{\sigma_B(t, a)}{\sigma_{\text{ref}}(t, a) + \epsilon}. \quad (21.2)$$

The precise bounded normalization should be chosen empirically. A high value indicates that the temporal realization of the same instrument is similar across brokers; a low value indicates divergence.

22 Cross-Asset Coherence

Let $\mathbf{r}_t$ be the vector of normalized asset returns and let $\Sigma_t = \text{Cov}(\mathbf{r}_t)$ be the local covariance matrix. The first principal component explains

$$C_A(t) = \frac{\lambda_1(\Sigma_t)}{\text{tr}(\Sigma_t)}. \quad (22.1)$$

This is a benchmark measure of cross-asset coherence. The proposed geometric framework should be tested *against* this quantity rather than simply replacing it.

23 Joint Market Coherence

A simple composite statistic is

$$C_{AB}(t) = C_A(t)C_B(t). \quad (23.1)$$

A more rigorous implementation would estimate a joint latent coherence factor. The conceptual interpretation is that $C_{AB} \approx 1$ means the market is coherent across both assets and broker observations.

24 The Geometric Market State Vector


Define the state vector

$$\mathbf{S}_t = \left(D_T(t) \quad V_G(t) \quad A_G(t) \quad C_A(t) \quad C_B(t) \quad D_{ta}(t) \quad D_{tb}(t) \quad \tau(t) \right)^\top. \quad (24.1)$$

This vector is the central candidate representation of market state. It deliberately avoids reducing the market to a single scalar.

25 The Coherence–Deformation Phase Diagram

A particularly useful projection is $x = C_{AB}$ against $y = D_T$. This produces the qualitative phase structure below.



The upper-right region is especially important. It represents $C_{AB} \gg 0$ combined with $D_T \gg 0$ — many observations changing coherently at the same time. Such a state may represent a broad regime transition.

26 A Geometry-Driven Dynamic Price Channel

The framework can be used to construct a dynamic channel. Let m_t be a structural centerline. Define the geometric width

$$w_t = w_0 \exp(\alpha D_T(t) + \beta V_G(t) + \gamma \sigma_t). \quad (26.1)$$

The channel boundaries are $U_t = m_t + w_t$ and $L_t = m_t - w_t$. A more general asymmetric construction is

$$U_t = m_t + w_t^+, \quad L_t = m_t - w_t^-. \quad (26.2)$$

The widths may depend separately on the direction of deformation.

27 Beltrami Phase and Directional Asymmetry

The magnitude $|\mu|$ describes the amount of deformation. The phase $\arg(\mu)$ contains directional information in the complex representation. This suggests an extension

$$w_t^\pm = f^\pm(|\mu_t|, \arg(\mu_t), C_t). \quad (27.1)$$

The hypothesis is that the *direction* of geometric deformation may contain information that is discarded when only a scalar volatility measure is used. This must be empirically tested.

28 The Three Deformation Families

The framework distinguishes three primary deformations.

28.1 Temporal deformation

$$D_t = D_T(\Phi_t, \Phi_{t+\Delta t}).$$

28.2 Asset deformation

For assets a_i, a_j ,

$$D_a(a_i, a_j; t) = D_T(\Phi_{a_i, t}, \Phi_{a_j, t}). \quad (28.1)$$

28.3 Broker deformation

For brokers b_i, b_j ,

$$D_b(b_i, b_j; a, t) = D_T(\Phi_{a, b_i, t}, \Phi_{a, b_j, t}). \quad (28.2)$$

These quantities should not be conflated.

29 Mixed Deformation Statistics

The most informative statistics may be mixed. Define D_{ta} , D_{tb} , D_{ab} , whose discrete counterparts are based on $\Delta_t \Delta_a P$, $\Delta_t \Delta_b P$, and $\Delta_a \Delta_b P$.

The temporal-broker term D_{tb} is especially important because it distinguishes common market movement from observation-specific movement. The temporal-asset term D_{ta} distinguishes common factor movement from pair-specific movement. The asset-broker term D_{ab} tests whether apparent asset relationships are stable across observation sources.

30 A Discrete Quasiconformal Approximation via the Jacobian

A full complex quasiconformal model is computationally demanding. A practical first implementation can use the local Jacobian

$$J = \begin{pmatrix} f_x & f_y \\ g_x & g_y \end{pmatrix}.$$

Let its singular values be $s_1 \geq s_2 > 0$. Define

$$K_{\text{disc}} = \frac{s_1}{s_2}. \quad (30.1)$$

The corresponding deformation is

$$D_{\text{disc}} = \frac{1}{2} \log K_{\text{disc}}. \quad (30.2)$$

This gives a real-valued discrete analogue of quasiconformal dilatation. Only after establishing empirical value should the full Beltrami optimization be introduced.

31 Estimating a Financial Beltrami Field

Let $z = u + iv$ be a local market coordinate, $F_0(u, v)$ a reference market patch and $F_1(u, v)$ the observed patch. We seek $f : F_0 \rightarrow F_1$. The map may be estimated through

$$\hat{f} = \arg \min_f [\mathcal{L}_{\text{data}}(f) + \lambda \mathcal{R}_{\text{QC}}(f)]. \quad (31.1)$$

Here $\mathcal{L}_{\text{data}}$ measures reconstruction error and $\mathcal{R}_{\text{QC}}$ regularizes the deformation. The resulting

$$\hat{\mu} = \frac{\partial_{\bar{z}} \hat{f}}{\partial_z \hat{f}} \quad (31.2)$$

defines the empirical Beltrami field.

Admissibility and guard discipline

A market patch is admissible for this estimation only if explicit data-quality conditions hold:

1. synchronized timestamps;
2. stable instrument identity;
3. stable venue identity;
4. cross-rate consistency checks;
5. explicit missing-data treatment;
6. explicit price definition (bid, ask, mid, last, mark);
7. numerical validity of the quasiconformal estimate.

The last condition is a hard guard. If the estimated field violates $\|\hat{\mu}_t\|_\infty < 1$ on some region, no standard quasiconformal interpretation may be claimed on that domain — the estimator must flag and exclude it, not clip the value back into range. An admissibility failure is a result about the data, not a numerical nuisance to be smoothed away.

32 Surface Geometry and the Metric Tensor

For a broker-time surface $S(t, b)$, the induced metric is

$$g_{\alpha\beta} = \langle \partial_\alpha S, \partial_\beta S \rangle, \quad \alpha, \beta \in \{t, b\}. \quad (32.1)$$

The determinant $\det g = g_{tt}g_{bb} - g_{tb}^2$ measures local area scaling; the off-diagonal g_{tb} measures directional coupling; and the normalized form is $\rho_{tb} = g_{tb}/\sqrt{g_{tt}g_{bb}}$. The same construction can be applied to the (t, a) and (a, b) surfaces.

33 Market Geometry as a Fibered Structure

The tensor $P(t, a, b)$ can be interpreted as a family of broker fibers over the asset-time base:

$$\pi : (t, a, b) \mapsto (t, a). \quad (33.1)$$

For every (t, a) there exists a fiber $\mathcal{F}_{t,a}$. This suggests a fibered market geometry:

base = time–asset structure	fiber = broker realization.
-----------------------------	-----------------------------

The geometry of the fibers measures observation dispersion, while the geometry of the base measures market structure. The mixed derivative $\partial_t \partial_b P$ then measures how the fiber geometry evolves as the base moves through time.

34 Recovering a Latent Common Market Geometry

Suppose several brokers observe the same latent market process $X(t, a)$. Each broker produces

$$P(t, a, b) = \mathcal{O}_b[X(t, a)] + \eta_b(t, a), \quad (34.1)$$

where $\mathcal{O}_b$ is the broker observation operator and η_b is observation noise. The research problem becomes:

Can the common latent geometry X be recovered from multiple deformed observations?

This is a central conceptual bridge between multi-broker analysis and deformation geometry.

35 Broker Observation Operators

A broker may be modeled abstractly as

$$\mathcal{O}_b = \mathcal{A}_b \circ \mathcal{L}_b \circ \mathcal{T}_b, \quad (35.1)$$

where $\mathcal{T}_b$ is the timestamp transformation, $\mathcal{L}_b$ the liquidity and aggregation transformation, and $\mathcal{A}_b$ the quote construction. The observed price becomes $P_b = \mathcal{O}_b[X]$.

The objective is not necessarily to recover each operator exactly. Instead, the geometric framework attempts to identify the portion of the observed structure that is *invariant* across operators.

36 Cross-Broker Consensus Geometry

Let $\Phi_{b_1}, \dots, \Phi_{b_M}$ be broker-specific market patches. Define a consensus patch

$$\Phi^* = \arg \min_{\Phi} \sum_{b=1}^M D_T(\Phi, \Phi_b)^2. \quad (36.1)$$

This is analogous to a geometric Fréchet mean. Each broker then has a deformation distance $D_b = D_T(\Phi_b, \Phi^*)$. The consensus geometry provides a reference against which broker-specific deviations can be measured.

A linear alternative: the market compensator

For two admissible observation representations $\Phi^{(1)}, \Phi^{(2)}$ — for example two venue realizations of the same instrument — let J denote the involution exchanging them. Define the compensator

$$C_{\text{mkt}} = \frac{1}{2}(I + J), \quad (36.2)$$

so that $C_{\text{mkt}}(\Phi) = \frac{1}{2}(\Phi^{(1)} + \Phi^{(2)})$, and the antisymmetric residual

$$R_{\text{mkt}} = \frac{1}{2}(I - J)\Phi, \quad \Phi = C_{\text{mkt}}\Phi + R_{\text{mkt}}. \quad (36.3)$$

The symmetric component is a candidate common-market estimate — a linear counterpart of the Fréchet-type consensus patch (36.1) that requires no optimization. The residual records observation-specific structure.

Critically, $R_{\text{mkt}} = 0$ must never be inferred merely because the compensator has been *defined*. It must be measured, or proved under explicit assumptions.

37 Cross-Asset Consensus Geometry

The same construction can be applied across assets. Given $\Phi_{a_1}, \dots, \Phi_{a_N}$, define

$$\Phi_A^* = \arg \min_{\Phi} \sum_{a=1}^N D_T(\Phi, \Phi_a)^2. \quad (37.1)$$

The residual $D_a = D_T(\Phi_a, \Phi_A^*)$ measures how far an individual asset deviates from the common market geometry. This can be particularly useful for identifying relative strength and relative structural isolation.

38 Coordinated and Fragmented Regimes

The combination of coherence and deformation leads to four qualitative states.

Coherence	Deformation	Interpretation
High	Low	Stable coordinated regime
High	High	Coordinated transition
Low	Low	Fragmented / range regime
Low	High	Structural instability

The final state is particularly important. If $D_T \uparrow$ while $C_{AB} \downarrow$, then the market is becoming geometrically unstable without developing a coherent cross-market direction. This may correspond to noise, liquidity stress, or the early phase of a regime transition.

39 Research Hypotheses

Hypothesis 1 (Cross-broker coherence). *After correct temporal synchronization and normalization, liquid instruments exhibit significant cross-broker structural coherence relative to randomized broker assignments.*

Hypothesis 2 (Geometric transition). *The deformation statistic D_T increases systematically around independently identified market regime transitions.*

Hypothesis 3 (Incremental information). *The geometric state vector contains information not explained by standard volatility, correlation, PCA, cointegration, or technical indicators.*

Hypothesis 4 (Coordinated transition). *Periods with simultaneously high coherence and increasing geometric deformation are disproportionately associated with broad market transitions.*

Hypothesis 5 (Broker anomaly detection). *Broker-specific geometric deviations can identify feed-specific anomalies more effectively than simple absolute price differences.*

Hypothesis 6 (Triangular closure). *Changes in the geometry of the FX triangular residual τ precede or coincide with selected liquidity and volatility transitions.*

Hypothesis 7 (Coordinate robustness). *A properly defined geometric deformation score is more stable under admissible coordinate changes than raw-price indicators.*

For the closed-form torus estimator of Section 18.2, H7 holds by construction (GL(2)-invariance); as stated, the hypothesis remains open for every other deformation score.

40 Comparison with Correlation, PCA, Cointegration, and Optimal Transport

40.1 Correlation

Correlation measures linear dependence,

$$\rho_{ij} = \frac{\text{Cov}(r_i, r_j)}{\sigma_i \sigma_j}.$$

It does not directly measure local geometric deformation.

40.2 Principal component analysis

PCA solves $\Sigma v_k = \lambda_k v_k$. The proposed framework may use PCA coordinates but introduces additional deformation information. PCA should therefore be a benchmark, not ignored.

40.3 Cointegration

Cointegration identifies stable long-run combinations. The proposed framework can incorporate cointegration residuals as coordinates but does not reduce the market geometry to them.

40.4 Optimal transport

Optimal transport provides a powerful alternative way of comparing distributions of market states. It should be included as a benchmark in any serious empirical test.

40.5 Manifold learning

Diffusion maps, Isomap, UMAP, and related methods can recover nonlinear low-dimensional structure. The Teichmüller component is justified only if it provides incremental information or a theoretically useful deformation metric.

41 Null Models and Estimator Validation

A sophisticated geometric model is vulnerable to false discoveries. The following null models are therefore mandatory, together with two positive validation tests of the estimator itself.

41.1 Broker permutation

Randomly permute broker identities while preserving the temporal and asset marginal distributions. If the geometric signal remains unchanged, the broker dimension contains no meaningful information.

41.2 Time shuffle

Randomly permute temporal ordering within blocks. A temporal geometric signal should be destroyed.

41.3 Asset shuffle

Randomly permute asset identities while preserving marginal return distributions. This tests whether cross-asset geometry depends on actual market relationships.

41.4 Synthetic models

At minimum test:

1. independent random walks;
2. correlated Gaussian processes;
3. stochastic volatility;
4. jump-diffusion;
5. regime-switching processes;
6. latent-factor models with broker-specific observation noise.

41.5 Synthetic deformation recovery

Before any market data are touched, generate a known quasiconformal map with a prescribed Beltrami field and verify that the estimation pipeline of Section 31 recovers it. An estimator that cannot recover a known μ from clean synthetic data has no business interpreting noisy market data. This test is logically prior to every null model above: the null models ask whether the market contains structure; this test establishes whether the instrument can measure structure at all.

41.6 Coordinate-invariance test

The embedding is determined only up to admissible conformal transformations (Sections 16 and 19). Apply such transformations to the same data and verify that the intrinsic deformation score is unchanged. A score that varies under conformal coordinate changes measures the choice of embedding, not the market — and every downstream result inherits that defect.

42 Empirical Dataset: FX Core and Cross-Asset Extension

A first empirical experiment should contain at least $M \geq 3$ independent broker or price-feed sources and $N \geq 10$ liquid instruments. A core FX universe may include

$$\begin{aligned} &\text{EUR/USD, GBP/USD, USD/JPY, USD/CHF, AUD/USD,} \\ &\text{USD/CAD, NZD/USD, EUR/JPY, EUR/GBP, GBP/JPY.} \end{aligned}$$

This universe should be extended by liquid instruments from outside FX that are observed through the same venues — for example XAU/USD, a major equity-index CFD, and BTC/USD — so that the asset coordinate spans several asset classes and cross-asset findings are demonstrably not an artifact of currency structure alone. The currency-network machinery of Sections 6–8 then applies to the FX subgraph, while every instrument participates in the temporal and broker dimensions.

Five-minute bars provide an accessible initial test. Tick-level data should be used for a second-stage microstructure analysis.

43 Statistical Power: What the Dataset Can Decide

A null result reported without a power calculation is a non-statement: it cannot distinguish “there is no effect” from “the data could not have shown an effect of this size.” Because the framework fixes its hypotheses before measurement, the same discipline is applied to detectability. This section computes, before any data are collected, which of the hypotheses H1–H7 the dataset of Section 42 can decide.

Throughout, the significance level is $\alpha = 0.05$ (two-sided) and the target power is 0.80, hence

$$k = z_{0.975} + z_{0.80} \approx 2.80, \tag{43.1}$$

and the smallest detectable standardized effect from n independent observations is

$$d_{\min} = \frac{k}{\sqrt{n}}. \tag{43.2}$$

43.1 Three classes of hypotheses, three limiting sample units

The binding constraint is never the raw bar count. Two years of five-minute bars over thirteen instruments and three venues contain roughly 5.7×10^6 bar observations, but adjacent bars and overlapping estimation windows are strongly dependent. What limits each hypothesis is its own unit of independent information.

Class A — window statistics (H1, H3, H5, H7). The limiting unit is the effectively independent estimation window. Treating one trading day as the scale of effective independence gives roughly 250 windows per instrument-year; two years over ten or more instruments yield n in the low thousands even after discounting cross-sectional dependence. By (43.2), standardized effects of $d \approx 0.04$ – 0.18 are detectable — a tenth of a standard deviation. H1 additionally admits an exact permutation null (the broker permutation of the null-model section), which does not rely on distributional assumptions at all. H7 is decidable with any amount of data, since coordinate robustness is a property of the estimator, testable analytically and on synthetic maps. *Verdict: decidable on the initial dataset.*

Class B — event statistics (H2, H4, H6). The limiting unit is the independent regime transition, and transitions are rare. With n independent events,

n events	10	20	30	50	100	200
$d_{\min}$	0.89	0.63	0.51	0.40	0.28	0.20

Under the working assumption of 5–15 independent transitions per year and market (the event definition must be preregistered together with its expected count), a large effect ($d = 0.5$, requiring 31 events) becomes decidable after roughly 2–6 years on a single market; a moderate effect ($d = 0.3$, requiring 87 events) after 6–17 years. Pooling across markets multiplies the event count only to the extent that the markets are genuinely independent, which the next paragraph bounds. *Verdict: decidable in the first cycle only for large effects; moderate effects require multi-year accumulation or genuinely independent markets.*

Class C — economic significance (Stage VI). The estimation error of a Sharpe ratio is approximately $\text{SE}(S) \approx \sqrt{(1 + S^2/2)/T}$ with T in years [17]. The minimum detectable Sharpe ratio after T years is

$$S_{\min}(T) = \frac{k}{\sqrt{T - k^2/2}}, \quad (43.3)$$

which is undefined below $k^2/2 \approx 3.9$ years: on shorter samples *no* Sharpe ratio of any size is statistically distinguishable from zero.

T (years)	5	10	15	20
$S_{\min}$ per series	2.70	1.14	0.84	0.70

A portfolio of BR independent bets improves this by $\sqrt{\text{BR}}$ — but ten USD-quoted FX pairs are far fewer than ten independent bets, because a common dollar factor dominates bilateral USD rates [24]. An effective breadth of $\text{BR} \approx 2$ – 4 is a realistic assumption for the core universe, and it is precisely to raise this number that Section 42 extends the universe beyond FX. At a realistic structural-edge Sharpe of 0.25 per instrument, the required sample is 35 years ($\text{BR} = 4$) to 67 years ($\text{BR} = 2$). *Verdict: Stage VI cannot be decided on the initial dataset. The decidable content of the first empirical cycle is Stages I–V, and a Stage-VI null result will not be reported as evidence of absence.*

43.2 Direction of the approximations

Every approximation above errs toward optimism: overlapping windows and autocorrelation shrink the effective n ; the effective-breadth argument replaces a heterogeneous correlation matrix by an equal-correlation summary; and analytic Sharpe standard errors understate the true error under autocorrelated returns, so wherever a permutation or control distribution is available it takes precedence over (43.3). The detectabilities reported here are therefore upper bounds — the honest reading is “at most this much can be decided,” never “at least.”

44 Rolling Window Estimation

For every time t , define $W_t = [t - L, t]$. All geometric quantities must be estimated using information in W_t only. This includes:

- normalization parameters;
- reference geometries;
- PCA coordinates;
- deformation parameters;
- latent factors;
- model hyperparameters.

The resulting state $\mathbf{S}_t$ must therefore be measurable with respect to information available at time t .

45 Data Leakage and Look-Ahead Bias

The framework is especially vulnerable to look-ahead bias because geometric normalization may inadvertently use future observations. For a valid backtest,

$$\theta_t = F(P_s : s \leq t) \quad (45.1)$$

must hold for every estimated parameter θ_t . The future return $r_{t+\Delta}$ may only be used for evaluation. No future data may influence Φ_t , μ_t , K_t , C_t , or D_t .

46 Trading Interpretation: From Geometry to Regime

The proposed geometry should *not* directly predict price direction. A safer hierarchy is

$$\boxed{\text{geometry} \rightarrow \text{regime} \rightarrow \text{conditional strategy} \rightarrow \text{risk management}}$$

For example, $D_T \downarrow$ together with $C_{AB} \uparrow$ could indicate a stable coordinated regime. A transition state could satisfy $D_T \uparrow$, $C_{AB} \uparrow$. A fragmented state could satisfy $D_T \uparrow$, $C_{AB} \downarrow$.

Direction would then be supplied by a separate directional model. This prevents the geometric framework from being forced to answer a question it was not designed to answer.

47 Transaction Costs and Broker Reality

Broker-level differences must be interpreted carefully. A deviation ΔP is not economically exploitable unless

$$|\Delta P| > \text{spread} + \text{commission} + \text{slippage} + \text{latency costs} + \text{other execution costs}.$$

Therefore:

$$\boxed{\text{geometric broker divergence} \neq \text{arbitrage opportunity}}$$

The primary purpose of the broker dimension is structural measurement.

48 Feed Anomaly Detection

A useful secondary application is detection of feed-specific anomalies. Suppose B_1, B_2, B_3 observe approximately the same market movement but B_1 deviates strongly. Define

$$A_b(t, a) = D_T(\Phi_{a,b,t}, \Phi_{a,t}^*). \quad (48.1)$$

A large A_b conditional on high cross-broker consensus may indicate a broker-specific anomaly. This can be used as a data-quality layer before any trading model.

49 Market Regime Classification

The geometric state vector can be used in a classifier $R_t = f(\mathbf{S}_t)$. Possible regimes include

$$R_t \in \{\text{trend, range, transition, fragmentation, event, microstructure stress}\}.$$

The classifier should preferably be trained using independent labels — volatility regimes, structural breaks, or event windows — rather than defining labels from the same geometric variables.

50 Multi-Timeframe Extension

Introduce a scale coordinate $\ell \in \{\text{tick, 1s, 1m, 5m, 1h, 1d}\}$. The market tensor becomes

$$P(t, a, b, \ell). \quad (50.1)$$

The deformation field is now $D(t, a, b, \ell)$. A particularly interesting question is whether deformation propagates across time scales, $D_{\ell_1} \rightarrow D_{\ell_2} \rightarrow D_{\ell_3}$. A structural transition that appears simultaneously across several scales may be qualitatively different from a purely local microstructure event.

51 Order-Book and Liquidity-Depth Extension

If depth-of-market information is available, introduce liquidity depth q . The market tensor becomes

$$P(t, a, b, q). \quad (51.1)$$

The q dimension represents the order-book layer. The full structure could then distinguish *price geometry* from *liquidity geometry*. A transition in price geometry without corresponding liquidity deformation may have a different interpretation from a transition accompanied by severe order-book deformation.

52 Generalization Beyond FX

The framework is not inherently restricted to foreign exchange. The same mathematical architecture can be applied to:

- equities across exchanges;
- index futures across venues;
- commodities across trading venues;
- cryptocurrency markets across exchanges;

- options surfaces;
- yield curves;
- volatility surfaces.

The observation coordinate b can represent a broker, an exchange, a venue, or a data vendor. The asset coordinate can represent an instrument, a maturity, or a strike, depending on the market.

53 The Geometric Hierarchy: From Price Chart to Market Tensor

The framework can be summarized by the hierarchy $P(t)$ for an isolated chart, $P(t, a)$ for a multi-asset market, and

$$P(t, a, b)$$

for a multi-asset, multi-observer market. The final object is not merely a collection of charts. It contains directional derivatives P_t, P_a, P_b , mixed derivatives P_{ta}, P_{tb}, P_{ab} , and local metric structure.

The Teichmüller layer adds a deformation comparison $\Phi_t \rightarrow \Phi_{t+\Delta t}$. The resulting quantity $D_T = \frac{1}{2} \log K$ is intended to measure structural deformation. The essential distinction is therefore

$$\text{price movement} \neq \text{geometric movement}$$

and

$$\text{broker dispersion} \neq \text{temporal movement.}$$

This distinction is the foundation of the framework.

54 Research Programme

The proposed research programme consists of six stages.

Stage I — Data geometry

Construct $P(t, a, b)$ with rigorous synchronization.

Stage II — Classical geometry

Estimate P_t, P_a, P_b and the mixed terms.

Stage III — Surface geometry

Construct the surfaces $(t, a), (t, b), (a, b)$ and estimate their local metrics.

Stage IV — Quasiconformal geometry

Construct normalized market patches and estimate μ, K, D_T . On two-dimensional patches the extremal distortion has the closed form of Section 18.2; the optimization is only needed beyond dimension two.

Stage V — Statistical validation

Compare against volatility, correlation, PCA, cointegration, and technical indicators.

Stage VI — Economic validation

Evaluate out-of-sample regime classification, conditional returns, and risk-adjusted trading performance. Section 43 shows that this stage is not decidable on the initial dataset; it becomes meaningful only as data accumulate over years.

55 Falsification Criteria

The framework should be considered unsuccessful if:

1. geometric statistics are completely explained by volatility;
2. broker geometry disappears after simple spread adjustment;
3. cross-asset geometry adds no information beyond PCA;
4. deformation has no out-of-sample stability;
5. performance disappears after transaction costs;
6. results depend critically on one broker;
7. results depend critically on one arbitrary asset ordering;
8. the apparent signal disappears under reasonable synchronization perturbations.

The principal failure modes behind these criteria are embedding dependence, timestamp artifacts, liquidity-source changes, hyperparameter overfitting, and the false claim that a numerical Beltrami-like quantity is an exact Teichmüller distance. These are design constraints, not footnotes.

These conditions are important because mathematical complexity can create an illusion of explanatory power.

56 Implementation Architecture

A modular implementation can be organized as

```
data/  
  raw/  
  synchronized/  
  normalized/  
  
geometry/  
  tensor.py  
  currency_graph.py  
  surfaces.py  
  derivatives.py  
  metrics.py  
  
teichmueller/  
  patches.py  
  beltrami.py  
  distortion.py  
  distance.py
```

```

statistics/
    coherence.py
    regime.py
    null_models.py

```

```

backtest/
    signals.py
    execution.py
    costs.py
    evaluation.py

```

The separation is intentional. The data layer should not contain trading assumptions. The geometry layer should not contain performance optimization. The trading layer should consume validated geometric states.

57 Minimal First Experiment

A first test should remain deliberately small. Use $M = 3$ or more brokers and approximately $N = 10$ liquid instruments, FX pairs plus a small number of non-FX instruments. Construct synchronized five-minute bars. For every rolling window:

1. normalize all prices;
2. calculate log returns;
3. estimate broker dispersion;
4. estimate cross-asset covariance;
5. calculate triangular residuals on the FX subgraph;
6. calculate mixed derivatives;
7. construct local surfaces;
8. calculate discrete distortion;
9. calculate coherence;
10. compare the resulting state against subsequent market behavior.

Only after this basic experiment demonstrates an effect should the full quasiconformal optimization be introduced.

58 First Empirical Cycle: Implementation Results

Stages I and II of the research programme have been implemented and run once, on two years of synchronized one-minute bars (2024-08-20 to 2026-08-18) for twelve instruments observed through three cTrader venues. Every claim below has now passed an adversarial review whose explicit task was refutation: the Stage-I and Stage-II claims in a first round, and the measurement-floor and estimation-floor subsections in a second round (2026-08-31) whose reviewer independently reproduced every reported number from the database and re-derived the two key identities ($d_T = \frac{1}{2}d_{\text{hyp}}$, and the isotropy of the direction null across anisotropic populations) before the numbers were accepted. All numbers are reported together with what they do *not* show. One provenance restriction governs everything in this section: all feeds are demo accounts of the respective venues, and whether a demo feed reproduces the venue's live price stream exactly is itself unverified — until it is, the results describe the demo feeds of three venues.

58.1 Stage I: the tensor is constructible, and the venue axis has measurable admission costs

Admissibility was defined as the intersection over venues: a minute belongs to the tensor only if all $M = 3$ venues quote it. For FX and gold this costs about 3% of minutes, leaving roughly 7×10^5 common minutes per instrument. One instrument breaks the naive reading of “the same price at the same time”: for BTC/USD, 53.4% of common minutes have *no* price level that all three venues saw simultaneously, driven by measured standing level offsets (about +4.7 and −21.8 USD against a $\sim 63,700$ USD level). Cross-venue price identity is an assumption to be checked per instrument, not a fact.

Two implementation lessons generalize. First, the integer price scale is a property of the *venue*, not the instrument — one venue quotes BTC/USD with three decimals, the others with two, and a raw cross-venue comparison was off by a factor of ten; it was caught only because a reported 100.00% disjointness looked impossible. Every per-venue constant must be treated as per-venue. Second, the price side (bid, ask, mid) must be pinned empirically before any geometry is computed; on these feeds the bar closes match the bid side at seven of eight instruments checked — on a deliberately small running sample ($n = 14$ minutes per instrument at the time of writing, one instrument within noise of the mid), so this is an indication under continued verification, not a settled fact.

58.2 Stage II: the venue axis fails as a geometric direction — and succeeds as a replication axis

The persistence of cross-venue differences was measured through pairwise venue differences of the log price *level* — and the level/return distinction matters: everything durably venue-specific lives in the level difference and cannot, by construction, appear in the return autocorrelation of any single venue that served as its control. The headline result of a first, naive pass — strong positive autocorrelation of the venue difference across all FX cells — did not survive the adversarial review: 14% of the minutes, around the daily rollover, carry 64–88% of the variance of the venue difference (the missing market-phase control); the three venue pairs satisfy $d_{13} \equiv d_{12} + d_{23}$ identically and are two degrees of freedom, not three; and the acceptance margin sat 59 null standard deviations above zero at $n \approx 3.5 \times 10^5$ (39 at the core-session $n \approx 1.5 \times 10^5$), so the criterion could not fail. A rerun restricted to the core session failed in 7 of 60 cells, among them the most liquid pair, EUR/USD, in three of its six cells and in *both* sample halves — under the symbols-and-halves rule of this framework that is a null result, and the surviving statement is deliberately modest: the venue difference is not white, an independent per-minute sampling offset does not explain it, and its variance is concentrated where liquidity is thinnest.

The structural conclusion is quantitative. With $\text{Var}(d)/\text{Var}(r) \approx 0.023$, a (time, venue) patch is intrinsically anisotropic by a factor of about 43:1; any dilatation statistic computed on it is a constant plus noise, and normalizing the axes divides out exactly the quantity of interest. *As a geometric direction, the venue axis does not carry.* It carries as a *replication axis*: M venues are M estimates of the same geometric object, and their spread is an error bar that requires no model, no bootstrap, and no asymptotics — precisely the quantity whose absence made the naive Stage-II criterion powerless.

58.3 The measurement floor of the conformal shape

Using the closed form (18.3), the conformal shape τ of a pair of instruments is the conformal class of their 2×2 return covariance per (pair, venue, window). Two numbers are then reported side by side: the *floor* — the mean pairwise d_T between the three venues within the same window — and the *movement* — the d_T between consecutive windows of the same venue. Their quotient answers, before any hypothesis is formulated, whether the shape moves through time by more than the instrument error of observing it. This is the inversion announced in Section 2: where

multi-venue econometrics removes cross-venue dispersion as fragmentation noise [6], it serves here as the error bar itself; the analytic sampling distribution of geodesic distances between correlation matrices [16] offers a model-based floor that a future analysis can lay beside this measured one.

Over all 45 pairs of ten FX instruments at daily windows the median quotient is 9.8 (range 4.8–19.1; a four-major pilot gave 14.3 and was not representative). The movement is nearly constant across pairs; the floor varies by $3.4\times$ and correlates with the anisotropy of the chosen plane ($r \approx 0.52$ on the log floor, ≈ -0.05 on the log movement) with a measured exponent of 0.29 against a naive prediction of 0.5 — anisotropy is one driver of the floor, not the driver, and this attribution rests on a single sample. Unlike the Stage-II statistic, the floor is *not* rollover-dominated: restricting to the core session leaves it essentially unchanged, which is the scale invariance of τ paying off. The numerical resolution of (18.3) sits near 10^{-8} , six orders of magnitude below the measured floor.

What the quotient is not: it is not a finding — a pure random walk also moves its shape; it carries no placebo, which belongs to a hypothesis, not to calibration; and it does not choose the plane, which remains a preregistered modeling decision. What it is: the first framework quantity measured with its own error bar, and a precondition — had the quotient been near one, every Stage-IV hypothesis about τ would have been undecidable regardless of sample size. One more thing it is not, established by the preregistered follow-up below: it is not the right denominator for questions about shape *dynamics*. Against the estimation floor of the window covariance, the same movement is a factor of five smaller.

58.4 The estimation floor: what the quotient of 9.8 is measured against

The venue floor answers one question — is the shape movement venue noise? — and a second, preregistered measurement answered the harder one: is it more than the sampling noise of estimating a 2×2 covariance from one window? The instrument was validated before it touched market data, against four conjunctive controls: a registered χ^2 prediction for the null distance ($2n d_T^2 \rightarrow \chi_2^2$, hence a median null distance of $\sqrt{\ln 2/n}$); the analytic sampling law of geodesic correlation distances of [16] as a code-independent cross-check; exact Wishart/GL(2) invariance of the null; and, for the direction statistic below, the isotropy of the null across anisotropic populations. All four passed.

The results re-anchor the headline number. The empirical split-half floor — two parity-balanced halves of the same window — runs at 2.0–2.6 times the i.i.d. floor, an upper bound on the dependence inflation of covariance estimation noise. The daily shape movement then sits at a median of ≈ 2.0 times this estimation floor (per-pair range 1.6–2.6), not 9.8: under the preregistered two-sided criterion, 43 of 45 pairs lie within ± 2 standard errors of the decision threshold and are recorded as *undecidable*; the single pair exceeding it is compatible with the multiple-comparison random rate (0.74 expected), and its identity flips when the split parity is changed — a de-facto perturbation test that the one “hit” did not survive. The direction of the daily step (median horizontal share 0.315 against a simulated isotropy null of 0.5) separates the data from the i.i.d. null but not from a structureless model with time-varying per-instrument volatility, which reproduces the measured share; the registered all-cells criterion failed in 19 of 90 cells.

One statement survives with its limits attached, and one has had to be narrowed. The shape autocorrelation is non-diffusive in the measured band: over 120 daily lags the distance grows by a factor of 1.9 against a memoryless reference of ≈ 9 –10, and a noise-corrected diffusive extrapolation of the one-day step overshoots the measured 120-day value by a factor of 4.9. The escape route from that conclusion was named rather than hidden — and running it down produced both a correction and a measurement.

The escape route, correctly bounded. Earlier versions priced the escape at a floor of ≈ 0.106 , about 2.1 times the measured split-half floor. That figure is not an admissible solution: the model $d(k)^2 = F^2 + s^2 k$ forces $F \leq d(1)$, and $d(1) = 0.0994$ lies below it. It is the same boundary defect that the preregistration's own review had removed at the *lower* edge, reappearing unnoticed at the upper one. Constrained to the admissible range, the published lag table is fitted by $F = 0.0994$, $s = 0.0166$ at 6.8% relative error — so on that table the diffusive reading is not excluded at all, merely bounded at ≤ 1.95 times the floor rather than 2.1. This number is stated because it *weakens* the claim of this section; publishing only the convenient half would have been the wrong choice.

The volatility-clustering null, now measured. The GARCH-class null that decides the escape was preregistered and frozen before any measurement — a constant-shape null that keeps the observed per-minute radial path and bootstraps the direction marginals, model-free rather than a fitted GARCH(1,1), so that seasonality and news clustering enter by construction, with its own instrument validation, plausibility guards, and the consequence of every outcome fixed in advance — and it has since been run. All four instrument controls passed before market data were touched, both guards held, and all 90 registered cells were populated. The path-unshared floor came out at 0.049 and 0.053 in the two time halves, that is 0.98 to 1.00 times the split-half floor: the downward bias that the escape route required does not appear, and by the registered criteria the outcome is closure in every one of the 90 cells. Those cells are dependent units drawn from a single asset class, so the unanimity is one piece of evidence rather than ninety.

That closure covers less than its name suggests, and the limit is structural rather than statistical. Because d_T is scale-invariant, the suspected bias cannot act through the volatility *level* at all; it can act only through the day-to-day dispersion of the within-window fourth moment. No scenario we were able to construct pushes that channel beyond ≈ 1.4 , while the criterion required 1.68–2.01 — with two caveats that keep this from being a statement about all possible states. The bound is an upper limit over a handful of constructed cases, not a proof; and the two figures live in different aggregation conventions (median for the channel, root mean square for the criterion), a comparison the preregistration explicitly forbids, with the conversion unchecked and running in the direction that would narrow the gap. More importantly, a shape offset that is constant within the day and independent across days — carrying no shape dynamics whatever — reproduces the step size, the directional share and both floors, and both instruments are blind to it by construction. Fitted on the two half-curves with an admissible floor, that broader reading is excluded in one time half and not in the other. The correct verdict is therefore: *the estimation-noise version of the escape is closed; the broader version is undecided*, and the non-diffusive reading of the lag table cannot claim this run as support.

What the run yielded instead. Two by-products outlast the verdict. The three validation anchors turn out to be special cases of one closed form,

$$F = \sqrt{\ln 2/n} \cdot \sqrt{\kappa/2}, \quad \kappa = \frac{\mathbb{E} m^4}{(\mathbb{E} m^2)^2}, \quad (58.1)$$

with m the whitened radial part: the estimation floor of any shape measurement depends on the window length and one fourth moment, to within a few per cent. It is predictable without simulation, and the run was in effect a measurement of κ (≈ 9.6 for foreign-exchange minutes on daily windows). The companion column with isotropic directions differs by at most one per cent, so the non-Gaussian direction distribution contributes nothing to the floor. The methodological lesson is the sharper of the two: the reachable range of a null model should be computed *before* it is registered. Here such a calculation would have raised the question in advance — though not settled it: measured against the expectation actually registered (0.09), the required factor would have been 1.13, comfortably inside the simulated range, and only against the value the criterion

turned out to demand does it become 1.68–2.01. Whether one of the two registered outcomes was genuinely unreachable therefore hinges on the same convention question this run did not settle. The general point survives that ambiguity: a criterion can be preregistered with perfect discipline and still have a reachable range nobody computed.

The honest summary of shape dynamics at daily windows is therefore unchanged in substance and better bounded in reach: clearly above venue noise, barely above estimation noise, direction indistinguishable from unstructured volatility, and not explicable by estimation noise under a shared radial path — while a memoryless daily offset remains an open alternative that the present instruments cannot see. Separating it requires the shape autocorrelation at sub-daily windows, an instrument the preregistration had explicitly set aside as uninformative and which this run promotes to the decisive one.

59 Expected Contributions

If validated, the framework could contribute four distinct ideas.

59.1 A multidimensional market representation

$P(t, a, b)$ provides a unified representation of time, asset, and observation dimensions.

59.2 Mixed-dimensional market statistics

The quantities P_{ta} , P_{tb} , P_{ab} provide information that ordinary one-dimensional indicators discard.

59.3 Geometric regime measurement

The deformation D_T provides a candidate measure of structural market change.

59.4 Observation-consensus geometry

Multiple brokers provide multiple realizations of the same market, allowing the construction of a consensus market geometry.

60 Limitations

Several limitations must remain explicit.

1. Brokers are not independent observations in a strict statistical sense.
2. Many of the relevant markets — FX, OTC and CFD instruments, crypto-assets — are decentralized, and there is no single canonical “true” price.
3. The definition of a market patch is not unique.
4. The mapping from financial data to a complex coordinate system is a modeling decision.
5. Teichmüller distance is mathematically well-defined only after the underlying objects and admissible transformations have been specified rigorously.
6. Financial time series are nonstationary.
7. Apparent geometric structure may be a repackaging of standard statistical dependence.
8. All feeds used in the first empirical cycle (Section 58) are demo accounts, and whether a demo feed reproduces the venue’s live price stream exactly is unverified; until it is, every cross-venue number in this paper describes the demo feeds of three venues.

These limitations make empirical falsification essential.

61 Conclusion

This paper proposes a structure-first framework for market analysis based on the multidimensional observation tensor

$$P(t, a, b)$$

where t is time, a the asset, and b the broker. The key conceptual step is to recognize that the three dimensions have different meanings. Movement along t is temporal market evolution; movement along a is cross-asset variation; movement along b is cross-observer variation. The mixed terms

$$P_{ta}, \quad P_{tb}, \quad P_{ab}$$

then describe how these dimensions interact.

The framework subsequently introduces local market patches and quasiconformal deformation. The Beltrami coefficient μ measures local deformation, while $K = (1 + |\mu|)/(1 - |\mu|)$ measures its magnitude. The proposed market deformation distance is

$$D_T = \frac{1}{2} \log K_{\min}.$$

The ultimate market state is represented not by one indicator but by a vector containing deformation, coherence, mixed-dimensional coupling, and consistency:

$$\mathbf{S}_t = (D_T, V_G, A_G, C_A, C_B, D_{ta}, D_{tb}, \tau).$$

The central empirical question is whether this state vector contains information about market regimes that cannot be reduced to standard measures such as volatility, correlation, PCA, cointegration, or conventional technical indicators.

The framework therefore makes a deliberately strong claim about the *research question*, but only a deliberately weak claim about the *result*:

$$\boxed{\text{The geometry is a hypothesis to be tested, not a result to be assumed.}}$$

If the empirical tests validate the construction, the next stage is not merely another indicator. It is the construction of a geometrically defined market state space in which trend, transition, fragmentation, and observation-specific distortion become different regions of the same mathematical object.

A Notation

Symbol	Meaning
$P(t, a, b)$	price of asset a at time t observed by broker b
$x(t, a, b)$	logarithmic price
r_t	temporal log return
t	time coordinate
a	asset / instrument coordinate (currency pair in the FX example)
b	broker / observation coordinate
$\mathcal{T}$	common time grid
$\mathcal{A}$	asset universe

$\mathcal{B}$	broker universe
$\mathcal{C}$	currency universe
B	currency graph incidence matrix
τ	triangular consistency residual
Φ_t	normalized market patch
μ	Beltrami coefficient
K	quasiconformal dilatation
D_T	geometric deformation distance
V_G	geometric velocity
A_G	geometric acceleration
C_A	cross-asset coherence
C_B	cross-broker coherence
D_{ta}	temporal-asset deformation
D_{tb}	temporal-broker deformation
D_{ab}	asset-broker deformation
C_{mkt}	market compensator
R_{mkt}	venue/asymmetry residual
$\mathcal{T}_{\text{mkt}}$	proposed market Teichmüller space

B Core Equations

The complete framework can be summarized by the following equations.

Market tensor

$$P = P(t, a, b).$$

Logarithmic coordinate

$$x = \log P.$$

First-order geometry

$$dP = P_t dt + P_a da + P_b db.$$

Second-order geometry

$$d^2P = P_{tt} dt^2 + P_{aa} da^2 + P_{bb} db^2 + 2P_{ta} dt da + 2P_{tb} dt db + 2P_{ab} da db.$$

Beltrami coefficient

$$\mu = \frac{f_{\bar{z}}}{f_z}.$$

Dilatation

$$K = \frac{1 + |\mu|}{1 - |\mu|}.$$

Geometric deformation

$$D_T = \frac{1}{2} \log K_{\min}.$$

Geometric velocity

$$V_G = \frac{D_T(\Phi_t, \Phi_{t+\Delta t})}{\Delta t}.$$

Cross-asset coherence

$$C_A = \frac{\lambda_1(\Sigma)}{\text{tr}(\Sigma)}.$$

Broker dispersion

$$\sigma_B^2 = \text{Var}_b(r_t).$$

Market state

$$\mathbf{S}_t = (D_T, V_G, A_G, C_A, C_B, D_{ta}, D_{tb}, \tau).$$

C Recommended Empirical Benchmark Matrix

Model	Input	Purpose
Baseline	Returns	Raw predictive benchmark
Volatility	Realized volatility	Regime benchmark
Correlation	Return covariance	Cross-asset benchmark
PCA	Factor concentration	Common-factor benchmark
Cointegration	Residuals	Relative-value benchmark
Technical	RSI / MA / channels	Trading benchmark
Geometry	D_T	Structural benchmark
Full model	$\mathbf{S}_t$	Proposed framework

D Final Research Criterion

The framework should ultimately be evaluated according to

Identification + Out-of-Sample Stability + Incremental Information
+ Economic Significance + Replicability.

Mathematical sophistication alone is insufficient. The central scientific test remains

$$I(\mathbf{S}_t; \text{future market state}) > I(\mathbf{B}_t; \text{future market state}), \quad (\text{D.1})$$

where $\mathbf{B}_t$ denotes an appropriate benchmark state vector. Only if the geometric state provides statistically *and* economically significant incremental information should the Teichmüller component be considered empirically meaningful.

E Research-Status Matrix

Every component of the framework carries an explicit status, so that established mathematics is never silently blended with untested proposals.

Component	Status	Requirement
Quasiconformal / Teichmüller mathematics	Established	Standard theory [1, 2, 13]
Geometry in finance	Established	Literature comparison [15, 20, 21]
Teichmüller concepts in financial time series	Prior art	Cite and delimit [14]
Market tensor $P(t, a, b)$	Proposed	Formalize and test
Multi-venue market surface	Proposed	Synchronized data
Beltrami market field	Proposed	Numerical validation (Section 41.5)
Teichmüller-type deformation score	Proposed	Synthetic + empirical validation
Market compensator	Hypothesis	Test as representation operator
Degeneracy of D_T on simply connected patches	Proved	Riemann mapping theorem (Section 18.1)
Closed-form d_T on torus moduli	Established	Teichmüller's theorem (Section 18.2)
H7 for the torus estimator	Theorem	GL(2) invariance (Section 18.2)
Venue axis as geometric direction	Refuted (first cycle)	Section 58
Venue axis as replication axis	Supported	Error-floor measurement (Section 58)
Shape movement above estimation noise	Undecidable	Preregistered split-half floor (Section 58)
Diffusive reading of the shape ACF	Partly excluded	Preregistered volatility-clustering null: estimation-noise version closed, memoryless daily offset undecided (Section 58)
Estimation floor as a closed form	Measured	Equation (58.1); the three validation anchors ($\kappa = 2, 1, 6.76$) are special cases of it, but two of them share one registered χ^2 prediction, and no derivation is given here
Statistical power of the initial dataset	Computed	Upper bounds (Section 43)
Trading edge	Unknown	Not decidable on the initial dataset (Section 43, Class C)

F Open Mathematical Questions

1. What is the correct intrinsic market surface associated with $P(t, a, b)$?
2. What equivalence relation should define a financial Teichmüller space?
3. Which transformations are coordinate changes, and which are economically meaningful deformations?
4. Can a discrete Beltrami field be defined directly on the market tensor, without an arbitrary two-dimensional embedding?

5. Is there a natural graph- or discrete-Teichmüller formulation for cross-venue market data?
6. Can FX triangular closure be incorporated as a geometric constraint rather than a separate residual statistic?
7. Is there a useful analogue of a mapping-class or Veech-group action for recurring market geometries?
8. Can the geometry distinguish information arrival from microstructure noise?

Question 4 is the sharpest: an affirmative answer would remove the embedding choice of Section 17 — and with it the largest single source of arbitrariness in the construction.

The first implementation cycle (Section 58) answers Question 2 for two-dimensional patches: the equivalence relation is conformal equivalence of flat tori, the resulting space is $\text{SPD}(2)/\mathbb{R}_+$, and the distance has the closed form (18.3). The question remains open beyond dimension two, where the Teichmüller identification no longer holds.

References

- [1] L. V. Ahlfors, *Lectures on Quasiconformal Mappings*, Van Nostrand, 1966.
- [2] K. Astala, T. Iwaniec, and G. Martin, *Elliptic Partial Differential Equations and Quasiconformal Mappings in the Plane*, Princeton University Press, 2009.
- [3] Bank for International Settlements, *Constrained liquidity provision in currency markets*, BIS Working Papers No. 1073, 2024. <https://www.bis.org/publ/work1073.htm>
- [4] Bank for International Settlements, *Triennial Central Bank Survey: Foreign Exchange Turnover*, BIS, various editions. <https://www.bis.org/statistics/rpfx.htm>
- [5] G. Bucchieri, S. Marmi, and R. N. Mantegna, Evolution of correlation structure of industrial indices of U.S. equity markets, *Physical Review E*, 88(1), 012806, 2013.
- [6] G. F. Dias and K. Schweikert, Integrated variance estimation for assets traded in multiple venues, *Journal of Econometrics*, 255, 106244, 2026. doi:10.1016/j.jeconom.2026.106244.
- [7] M. P. do Carmo, *Differential Geometry of Curves and Surfaces*, Prentice Hall, 1976.
- [8] Encyclopedia of Mathematics, *Quasi-conformal mapping*. https://encyclopediaofmath.org/wiki/Quasi-conformal_mapping
- [9] Encyclopedia of Mathematics, *Riemann surfaces, conformal classes of*. https://encyclopediaofmath.org/wiki/Riemann_surfaces,_conformal_classes_of
- [10] Encyclopedia of Mathematics, *Quadratic differential*. https://encyclopediaofmath.org/wiki/Quadratic_differential
- [11] E. F. Fama and K. R. French, Common risk factors in the returns on stocks and bonds, *Journal of Financial Economics*, 33(1), 3–56, 1993.
- [12] F. P. Gardiner and W. J. Harvey, Universal Teichmüller space, arXiv:math/0012168, 2000.
- [13] F. P. Gardiner and N. Lakic, *Quasiconformal Teichmüller Theory*, Mathematical Surveys and Monographs, Vol. 76, American Mathematical Society, 2000.
- [14] K. Kanjamapornkul, R. Pinčák, S. Chunitipaisan, and E. Bartoš, Support Spinor Machine, *Digital Signal Processing*, 70, 59–72, 2017. doi:10.1016/j.dsp.2017.07.023.

- [15] K. Kanjamapornkul, R. Pinčák, and E. Bartoš, Cohomology theory for financial time series, *Physica A: Statistical Mechanics and its Applications*, 546, 122212, 2020. doi:10.1016/j.physa.2019.122212.
- [16] A. Kuketayev, Connecting Riemannian geometry and statistical inference for correlation matrices, arXiv:2608.27209, 2026.
- [17] A. W. Lo, The statistics of Sharpe ratios, *Financial Analysts Journal*, 58(4), 36–52, 2002.
- [18] B. B. Mandelbrot, The variation of certain speculative prices, *The Journal of Business*, 36(4), 394–419, 1963.
- [19] M. C. Münnix, T. Shimada, R. Schäfer, F. Leyvraz, T. H. Seligman, T. Guhr, and H. E. Stanley, Identifying states of a financial market, *Scientific Reports*, 2, 644, 2012.
- [20] P. Papaioannou and A. N. Yannacopoulos, The Shape of Markets: Machine learning modeling and prediction using 2-manifold geometries, arXiv:2511.05030, 2025.
- [21] R. Pinčák and K. Kanjamapornkul, GARCH(1,1) model of the financial market with the Minkowski metric, *Zeitschrift für Naturforschung A*, 73(8), 669–684, 2018. doi:10.1515/zna-2018-0199.
- [22] H. S. Shin, *Risk and Liquidity*, Oxford University Press, 2010.
- [23] Y. Thanwerdas and X. Pennec, Geodesics and curvature of the quotient-affine metrics on full-rank correlation matrices, in *Geometric Science of Information (GSI 2021)*, Lecture Notes in Computer Science, Vol. 12829, 93–102, Springer, 2021.
- [24] A. Verdelhan, The share of systematic variation in bilateral exchange rates, *The Journal of Finance*, 73(1), 375–418, 2018.